

# Assignment 1 Report

Anay Kulkarni

## Part 1A

The goal of these experiments was to evaluate the performance of different configurations of neural networks (NNs) for sentiment analysis, using GloVe embeddings and a Deep Averaging Network (DAN) architecture. Each network was tested across 50 epochs, and the results are presented based on test and train accuracy. We explored variations in loss functions, activation functions, optimizers, and the depth of the neural network.

### Experimental Setup

- **Dataset:** Movie review sentences classified as positive (1) or negative (0).
- **Embedding:** GloVe embeddings with 300 dimensions, pre-trained.
- **Neural Network Architectures:**
  - 2 fully connected (FC) layers and 1 dropout layer
  - 3 fully connected layers and 2 dropout layers
- **Loss Functions:** Negative Log-Likelihood (NLL) and Cross-Entropy.
- **Activation Functions:** ReLU, Tanh, and Sigmoid.
- **Optimizers:** AdamW, Stochastic Gradient Descent (SGD) with momentum.
- **Batch Sizes:** 16 and 32.
- **Number of Epochs:** 50.

### Experiment 1: Effect of Activation Functions and Loss Functions




Figure 1: Training Accuracy




Figure 2: Test Accuracy

Table 1: Performance with Different Activation and Loss Functions

Epoch	ReLU + NLL (Train)	ReLU + NLL (Test)	Tanh+Sigmoid + NLL (Train)	Tanh+Sigmoid + NLL (Test)
5	0.820	0.790	0.823	0.787
10	0.896	0.811	0.901	0.812
15	0.936	0.810	0.942	0.814
20	0.958	0.804	0.964	0.815
30	0.985	0.796	0.981	0.811
50	0.998	0.795	0.990	0.800

### Discussion:

- **ReLU with NLL** performed better initially, with a higher training accuracy after 50 epochs.
- **Tanh+Sigmoid with NLL** achieved better generalization (test accuracy 0.815) but had slower initial progress compared to ReLU.
- The overfitting trend is more evident with ReLU, as the training accuracy approached 99.8%, while test accuracy remained stagnant or slightly decreased after 20 epochs.

---

## Experiment 2: Comparison Between NLL and Cross-Entropy Loss




Figure 3: Training Accuracy

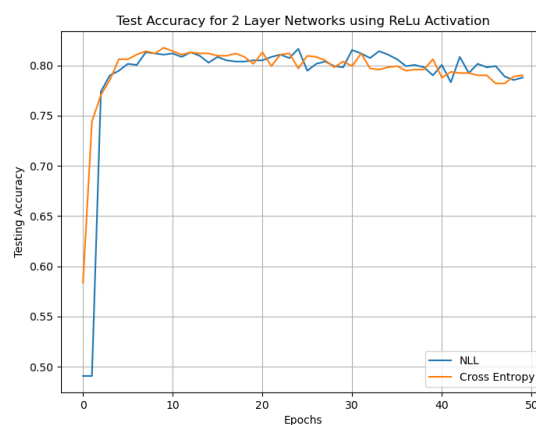


Figure 4: Test Accuracy

Table 2: NLL vs Cross-Entropy with ReLU Activation

Epoch	ReLU + NLL (Train)	ReLU + NLL (Test)	ReLU + Cross Entropy (Train)	ReLU + Cross Entropy (Test)
5	0.824	0.795	0.828	0.806
10	0.900	0.811	0.896	0.818
15	0.939	0.803	0.936	0.812
25	0.976	0.817	0.976	0.797
50	0.997	0.788	0.998	0.790

### Discussion:

- Cross-Entropy performed marginally better in terms of test accuracy during the initial epochs (e.g., 0.806 at epoch 5), suggesting it was able to generalize better early on.
- However, the difference between Cross-Entropy and NLL became negligible as the epochs increased. Both functions showed slight overfitting.

### Experiment 3: Effect of Depth in Neural Network

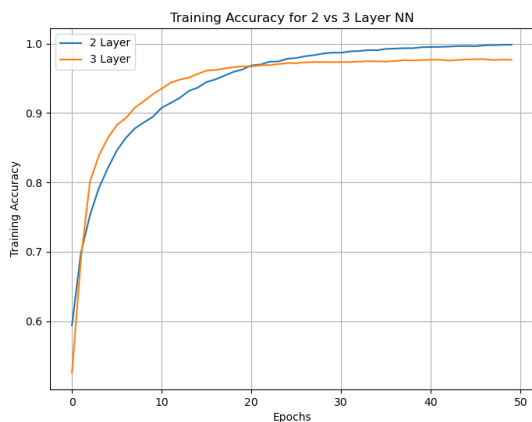


Figure 5: Training Accuracy




Figure 6: Test Accuracy

Table 3: 2-Layer vs. 3-Layer NN with Different Activation Functions

Epoch	2-Layer (ReLU)		3-Layer (Tanh+Sigmoid)	
	(Train)	(Test)	(Train)	(Test)
5	0.821	0.792	0.864	0.802
10	0.895	0.814	0.927	0.812
20	0.963	0.810	0.967	0.810
30	0.987	0.815	0.973	0.808
50	0.999	0.791	0.977	0.795

#### Discussion:

- The **3-layer** NN showed more robust generalization during the first 20 epochs, consistently outperforming the 2-layer NN in test accuracy, particularly in the early epochs.
- However, the difference between the two networks diminished as the number of epochs increased.

### Experiment 4: Hyperparameter Tuning (Batch Size and Optimizer)

Table 4: Effect of Batch Size and Optimizer on 3-Layer NN (with Tanh+Sigmoid)

Epoch	Batch Size: 16 + AdamW (Train)	Batch Size: 16 + AdamW (Test)	Batch Size: 32 + SGD (Train)	Batch Size: 32 + SGD (Test)
5	0.861	0.807	0.661	0.732
10	0.926	0.814	0.745	0.742
25	0.970	0.812	0.770	0.737
50	0.976	0.814	0.800	0.802

#### Discussion:




Figure 7: Training Accuracy




Figure 8: Test Accuracy

- **AdamW optimizer** with a smaller batch size of 16 performed better overall in terms of training and test accuracy.
- The **SGD optimizer** with momentum (0.9) and a batch size of 32 showed slower training progress but improved generalization after 50 epochs.

## Conclusion

- **Activation Functions:** The combination of Tanh and Sigmoid performed slightly better than ReLU in terms of generalization, though ReLU enabled faster training.
- **Loss Functions:** Cross-Entropy and NLL had similar performance, with Cross-Entropy showing a slight advantage in early epochs.
- **Network Depth:** The 3-layer NN provided better generalization initially but did not provide significant improvement over the 2-layer NN after prolonged training.
- **Optimizers:** AdamW with a smaller batch size yielded the best overall performance, with steady and consistent improvements in both train and test accuracy.
- **Overall:** The best configuration was the 3 layer fully connected Neural Network using Negative Log Likelihood Loss function and AdamW (Adaptive Moment Estimation) with Weight Decay.
  - Learning Rate: 0.0001
  - Weight Decay: 0.01
  - Hidden Units: 100
  - Batch Size: 16
  - Glove Dimensions: 300

## Additional Notes:

### Impact of (Over) Preprocessing on Data

- The dataset underwent additional preprocessing steps, including stemming, converting text to lowercase, and removing stop words, using the NLTK library. While these are common preprocessing methods, they actually led to a decline in accuracy in this context. The highest accuracy obtained by the models was 74% on the preprocessed data. This drop in performance could be due to the nature of the reviews, which typically contain short sentences. Aggressive preprocessing, such as stop-word removal and stemming, may have overly simplified the data, reducing its complexity and leading to weaker test performance.

## Part 1B

### Effects of Training Embeddings from Scratch vs. Fine-Tuning Pre-trained GloVe Embeddings on Model Convergence and Performance

The experiment evaluates the impact of two approaches to word embeddings—fine-tuning pre-trained GloVe embeddings and training random embeddings from scratch—on the performance and convergence of a 3-layer neural network for sentiment classification. The comparison focuses on how these embedding strategies influence training speed, final accuracy, and generalization to unseen test data.

#### Experimental Setup:

1. **Model Architecture:**
  - 3-layer Neural Network (NN) with an embedding layer, three fully connected layers, two dropout layers, and softmax layer for output
2. **Embedding Types:**
  - **GloVe Embeddings:** Pre-trained on large corpora and fine-tuned during training.
  - **Random Embeddings:** Initialized with random weights and trained from scratch.
3. **Loss Function:** Negative Log-Likelihood (NLL)
4. **Activation Functions:** TanH and Sigmoid
5. **Optimizer:** Adam with weight decay
6. **Evaluation Metric:** Train and test accuracy over 100 epochs.

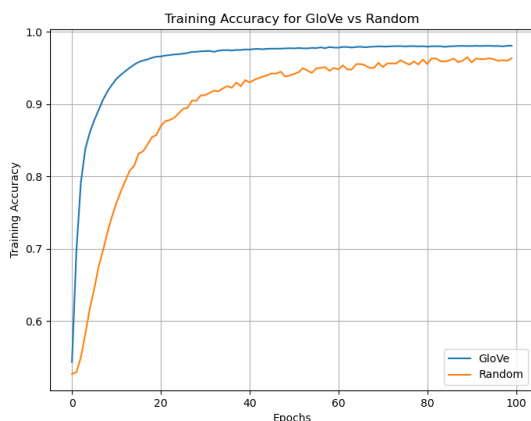


Figure 9: Training Accuracy

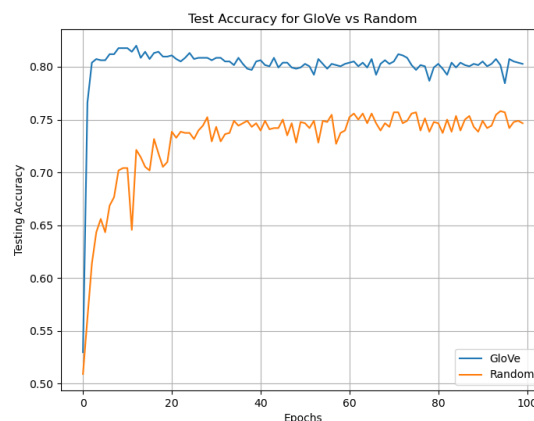


Figure 10: Test Accuracy

#### Results:

Epoch	Pre-trained GloVe (Train Accuracy)	Pre-trained GloVe (Test Accuracy)	Random Embeddings (Train Accuracy)	Random Embeddings (Test Accuracy)
5	0.861	0.806	0.617	0.656
10	0.927	0.818	0.745	0.704
15	0.955	0.814	0.815	0.705
20	0.966	0.810	0.857	0.710
25	0.969	0.813	0.887	0.737
30	0.973	0.806	0.912	0.729
50	0.977	0.799	0.940	0.748
75	0.980	0.802	0.961	0.756
100	0.981	0.803	0.964	0.747

## Analysis:

### 1. Convergence Speed:

- **GloVe Embeddings:** The model using fine-tuned GloVe embeddings reaches an 80% test accuracy within the first 5 epochs and achieves nearly 98% train accuracy by epoch 25. Test accuracy stabilizes around 80-81% early in training.
- **Random Embeddings:** Training from scratch results in much slower convergence. At epoch 5, the test accuracy is 65.6%, significantly lagging behind GloVe's 80.6%. It takes until epoch 75 for random embeddings to achieve similar accuracy (75.6%).

### 2. Final Accuracy:

- **GloVe Embeddings:** By epoch 100, the test accuracy fluctuates around 80.3%. This suggests strong generalization, as the model reaches a performance plateau early on, with minimal overfitting.
- **Random Embeddings:** After 100 epochs, the random embeddings achieve a test accuracy of 74.7%. This indicates lower generalization compared to the GloVe embeddings, which could be attributed to the lack of prior knowledge in the random vectors.

### 3. Generalization and Overfitting:

- **GloVe Embeddings:** The relatively close gap between training and test accuracies (around 18% by epoch 100) suggests that fine-tuning pre-trained GloVe embeddings helps maintain good generalization.
- **Random Embeddings:** There is a larger gap between training and test accuracy, especially in later epochs (96.4% train accuracy vs. 74.7% test accuracy at epoch 100), signaling overfitting as the model learns more detailed but potentially less useful information.

### 4. Training Time:

- The total training time for both models was similar (around 360 seconds), but **the GloVe model required fewer epochs to reach optimal test performance**. This indicates that GloVe embeddings, which come with pre-learned semantic relationships, allow for faster convergence despite similar computational costs.

**Conclusion:** Fine-tuning pre-trained GloVe embeddings significantly improves model convergence, resulting in faster achievement of high accuracy compared to training embeddings from scratch. Additionally, models trained with GloVe embeddings exhibit better generalization and lower overfitting, making them a preferable choice when available. On the other hand, training embeddings from scratch takes longer to converge and results in higher risk of overfitting, especially when training data is limited.

---

## PART 2A

In this experiment, we compared the performance of a Deep Averaging Network (DAN) for sentiment analysis using two different tokenization strategies: 1. **Subword Tokenization** with Byte Pair Encoding (BPE) 2. **Word-level Tokenization**

The goal was to understand the impact of using subword tokenization in terms of model performance, training time, and vocabulary compression.

### 1. Experimental Setup

- **Model:** Deep Averaging Network (DAN)
- **Embedding:** Random embeddings for both subword and word-level tokens
- **Training Duration:** 100 epochs
- **Dataset:** Movie reviews (vocabulary sizes provided below)
- **Metrics:** Train accuracy, test accuracy, training time, evaluation time

## 2. Results




Figure 11: Training Accuracy




Figure 12: Test Accuracy

### 2.1 Vocabulary Comparison

Metric	Word-Level Vocabulary	Subword-Level Vocabulary (Initial)	Subword-Level Vocabulary (After BPE)	Compression Ratio
<b>Training Data Vocabulary Size</b>	14,830	64	471	31.49
<b>Test Data Vocabulary Size</b>	4,339	53	464	9.35

### 2.2 Model Performance Comparison (Accuracy)

Epoch	Subword Tokenization (Train Accuracy)	Subword Tokenization (Test Accuracy)	Word-Level Tokenization (Train Accuracy)	Word-Level Tokenization (Test Accuracy)
5	0.538	0.509	0.616	0.640
10	0.646	0.514	0.752	0.717
15	0.683	0.503	0.832	0.717
20	0.702	0.497	0.871	0.742
25	0.719	0.494	0.895	0.745
30	0.731	0.500	0.915	0.755
50	0.752	0.507	0.946	0.758
75	0.756	0.508	0.958	0.745
100	0.771	0.513	0.962	0.763

### 2.3 Time Comparison

Task	Subword Tokenization	Word-Level Tokenization
<b>Training Time (sec)</b>	54.29	371.14
<b>Evaluation Time (sec)</b>	1.99	1.69

### 3. Analysis

#### 3.1 Vocabulary Compression

- **Subword tokenization** significantly reduced the vocabulary size, as shown by the compression ratios. In the training data, the subword-level vocabulary went from 64 initial tokens to 471 after applying BPE, compared to the word-level vocabulary size of 14,830.
- This compression effect was similarly observed in the test data, where the vocabulary was reduced from 53 to 464 subword tokens. The compression ratios were 31.49 for the training data and 9.35 for the test data.

#### 3.2 Performance Analysis

- **Word-level tokenization** consistently outperformed subword tokenization in terms of both **train** and **test accuracy**. By the 100th epoch, word-level tokenization reached a test accuracy of **0.763**, while subword tokenization achieved a lower test accuracy of **0.513**.
- While subword tokenization had an initially slower improvement in accuracy, the model still demonstrated some capability to generalize. However, it never caught up to word-level tokenization, indicating that more training or tuning may be necessary for subword embeddings to reach comparable performance.

#### 3.3 Training and Evaluation Time

- **Training with subword tokenization** was significantly faster, completing in **54.29 seconds**, compared to **371.14 seconds** for word-level tokenization. This time is summed over all 100 epochs. This speed-up is due to the smaller subword-level vocabulary size, which reduces the computational complexity during training and evaluation.

#### 3.4 Potential for Generalization

- Subword tokenization has an advantage in generalizing to unseen words, as it breaks down words into smaller units, which can help in low-resource scenarios or when the test data contains many out-of-vocabulary words.
- Despite this, the **reduced accuracy** with subword tokenization might indicate that the model requires further refinement, such as hyperparameter tuning or pretraining on a larger dataset, to leverage the strengths of subword representations.

### 4. Conclusion

The experiment showed that while **subword tokenization** leads to faster training and smaller vocabulary sizes, it lagged behind **word-level tokenization** in terms of accuracy. Subword-level approaches may require additional optimization or pretraining to improve performance, but they offer significant benefits in terms of generalization and speed.

---

## Effect of Vocabulary Size Using Subword Tokenization in Neural Networks

### Introduction

This experiment investigates the effect of vocabulary size on the performance of a neural network trained on sentiment analysis task, utilizing subword-level tokenization through Byte Pair Encoding (BPE). A three-layer neural network was employed for training. Training was conducted from scratch using random weights for embeddings.






Figure 13: Accuracy for Subword-Level Tokenization (BPE) as a function of vocabulary length

## Results

### 1. Vocabulary Size and Compression Ratios

The following table summarizes the vocabulary sizes and compression ratios before and after applying the BPE algorithm:

BPE Merges	Training Data - Before BPE	Training Data - After BPE	Test Data - Before BPE	Test Data - After BPE	Compression Ratio (Training)	Compression Ratio (Test)
100	64	136	53	126	109.04	34.44
200	64	213	53	205	69.62	21.17
300	64	296	53	293	50.10	14.81
400	64	385	53	375	38.52	11.57
500	64	471	53	464	31.49	9.35
600	64	561	53	553	26.43	7.85
700	64	651	53	637	22.78	6.81
800	64	742	53	726	19.99	5.98
900	64	832	53	815	17.82	5.32
1000	64	923	53	899	16.07	4.83

### 2. Model Accuracy Across Epochs

Each NN is trained for 100 epochs over increasing vocabulary sizes. The peak accuracy achieved during the 100 epochs at a given vocabulary size is recorded. The following table presents the results for each vocabulary:

Vocabulary Size	Training Time (s)	Evaluation Time (s)	Peak Train Accuracy	Peak Test Accuracy
136	48.52	1.97	0.648	0.523
213	46.32	1.90	0.700	0.540
296	52.58	2.13	0.733	0.541
385	44.25	1.73	0.754	0.557
471	47.71	1.73	0.771	0.540

Vocabulary Size	Training Time (s)	Evaluation Time (s)	Peak Train Accuracy	Peak Test Accuracy
561	53.18	1.88	0.788	0.522
651	51.05	1.68	0.803	0.517
742	55.08	1.81	0.814	0.546
832	48.71	1.60	0.828	0.515
923	52.04	1.74	0.842	0.547

## 2.1. Analysis

- **Accuracy Trends:**
  - **Training Accuracy:** There is a steady increase in peak training accuracy as the vocabulary size increases. The model with a vocabulary size of 923 achieved the highest training accuracy (0.842).
  - **Test Accuracy:** Test accuracy increases with increase in vocabulary size initially, with the best test accuracy observed at 385 (0.557). Then, the performance fluctuates and declines as vocabulary size increases further. However, a sharp increase is noted again at 923.
- **Training and Evaluation Times:**
  - Training time shows a general increase with larger vocabulary sizes, peaking at 55.08 seconds for a vocabulary size of 742. The evaluation time also tends to increase with larger vocabularies but remains within a close range across different sizes (1.6-2.1 seconds).

## 3. Conclusion

- **Accuracy:** While larger vocabularies generally led to better training accuracy, some fluctuations were observed in test accuracies, suggesting possible overfitting without enough generalization capabilities. A moderate vocabulary size of 385 showed the best test accuracy (0.557). Further exploration into fine-tuning and techniques to reduce overfitting might benefit overall performance
- **Computational Costs:** A clear trend can be observed that shows higher vocabulary compression is associated with shorter training times.

# PART 3

## Q1)

In the skip-gram model, given a center word, the objective is to predict the surrounding context words within a window size  $k$ . Here,  $k = 1$ , meaning the model considers one word to the left and one word to the right as context words.

Sentence	Center Word	Context Word	Training Pair
"the dog"	the	dog	(the, dog)
"the dog"	dog	the	(dog, the)
"the cat"	the	cat	(the, cat)
"the cat"	cat	the	(cat, the)
"a dog"	a	dog	(a, dog)
"a dog"	dog	a	(dog, a)

## (3a) Optimal Probabilities for the Training Set

Given a dataset  $D$  consisting of word-context pairs  $(x, y)$ , where  $x$  is the center word and  $y$  is the context word, the goal is to maximize the likelihood of the observed data under the skip-gram model. The likelihood  $L(\theta)$  of the observed data  $D$  under the model is the product of the probabilities of all word-context pairs in

the dataset. If we have  $N$  word-context pairs, this is expressed as:

$$L(\theta) = \prod_{(x,y) \in D} P(y | x; \theta)$$

Where:

- $P(y | x; \theta)$  is the probability of the context word  $y$  given the center word  $x$ , parameterized by  $\theta$  (the model parameters, which include the word and context embeddings).
- $D$  is the set of observed word-context pairs in the training data.

To simplify the optimization process, we usually work with the **log-likelihood** instead of the likelihood itself. Taking the logarithm of the likelihood function turns the product into a sum:

$$\log L(\theta) = \sum_{(x,y) \in D} \log P(y | x; \theta)$$

The objective is to find the parameters  $\theta$  that maximize this log-likelihood. We know that the likelihood is maximized by having the model match the empirical distribution of outcomes observed in the data.

Let  $P_{\text{emp}}(y | x)$  be the empirical probability of seeing context word  $y$  given center word  $x$ . This empirical distribution is based on the relative frequency of observing each pair  $(x, y)$  in the data:

$$P_{\text{emp}}(y | x) = \frac{\text{Count}(x, y)}{\text{Count}(x)}$$

Where:

- $\text{Count}(x, y)$  is the number of times the pair  $(x, y)$  occurs in the data.
- $\text{Count}(x)$  is the number of times the center word  $x$  appears in the data.

The training data includes two context words for "the" ("dog" and "cat"). Therefore,

$$P_{\text{emp}}(\text{dog} | \text{the}) = \frac{\text{Count}(\text{the}, \text{dog})}{\text{Count}(\text{the})} = \frac{1}{2}$$

$$P_{\text{emp}}(\text{cat} | \text{the}) = \frac{\text{Count}(\text{the}, \text{cat})}{\text{Count}(\text{the})} = \frac{1}{2}$$

### (3b) Nearly Optimal Word Vector for "the"

We have,

$$P(y | x) = \frac{e^{v_x \cdot c_y}}{\sum_{y'} e^{v_x \cdot c_{y'}}$$

We want to calculate  $\mathbf{v}_{\text{the}}$ , such that,

$$P(y | \text{the}) = \frac{e^{v_{\text{the}} \cdot c_y}}{\sum_{y'} e^{v_{\text{the}} \cdot c_{y'}}$$

$$P(y | \text{the}) = \frac{e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_y}}{e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{cat}}} + e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}}}}$$

To find a nearly optimal word vector for “the,” we need to achieve probabilities close to  $P(\text{dog} \mid \text{the}) = \frac{1}{2}$  and  $P(\text{cat} \mid \text{the}) = \frac{1}{2}$  using the softmax function.

$$P(\text{dog} \mid \text{the}) = \frac{e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}}}}{e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{cat}}} + e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}}}}$$

$$\frac{1}{2} = \frac{e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}}}}{e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{cat}}} + e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}}}}$$

Solving this equation, we get,

$$e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{cat}}} = e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}}}$$

i.e., a trivial solution that satisfies the above equation can be obtained by simply setting the dot products to zero.

$$\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}} \approx \mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{cat}} \approx 0$$

We know, the context embedding vectors for “dog” and “cat” are:

$$\mathbf{c}_{\text{dog}} = \mathbf{c}_{\text{cat}} = (0, 1)$$

Let,  $\mathbf{v}_{\text{the}} = (a, b)$

$$\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}} = (a, b) \cdot (0, 1)^T = 0$$

i.e., any vector of the form,  $\mathbf{v}_{\text{the}} = (a, 0)$ , where  $aa$  is a real number, could be an optimal vector for the word “the” in our dataset.

However, in practice, the value of the vectors would be randomly initialized, and the optimization algorithm will converge the weights of  $\mathbf{v}_{\text{the}}$  to a value that approximately takes the form ( $\text{weight}_{\text{random}}, 0$ )

---

### (3c) Training Examples Derived from These Sentences

Sentence	Center Word	Context Word	Training Pair
“the dog”	the	dog	(the, dog)
“the dog”	dog	the	(dog, the)
“the cat”	the	cat	(the, cat)
“the cat”	cat	the	(cat, the)
“a dog”	a	dog	(a, dog)
“a dog”	dog	a	(dog, a)
“a cat”	a	cat	(a, cat)
“a cat”	cat	a	(cat, a)

### 3d: A Set of Both Word and Context Vectors that Nearly Optimizes the Skip-Gram Objective

Let’s calculate the probabilities  $P(y \mid x)$  using the formula:

$$P(y \mid x) = \frac{\text{count}(x, y)}{\text{count}(x)}$$

Where:

- $\text{count}(x, y)$  is the number of times  $y$  appears as a context word for  $x$
- $\text{count}(x)$  is the total number of times  $x$  appears as a center word across all training pairs.

**Table of Probabilities:**

Probability $P(y x)$	Count(x, y)	Count(x)	Probability Value
$P(\text{dog} \text{the})$	1	2	0.5
$P(\text{cat} \text{the})$	1	2	0.5
$P(\text{the} \text{dog})$	1	2	0.5
$P(a \text{dog})$	1	2	0.5
$P(\text{the} \text{cat})$	1	2	0.5
$P(a \text{cat})$	1	2	0.5
$P(\text{dog} a)$	1	2	0.5
$P(\text{cat} a)$	1	2	0.5

Assume some trivial vector representations for center words,

$$\mathbf{v}_{\text{the}} = (0, 1)$$

$$\mathbf{v}_{\text{dog}} = (1, 0)$$

$$\mathbf{v}_{\text{cat}} = (1, 0)$$

$$\mathbf{v}_a = (0, 1)$$

Similar to the solution from (3b), we can calculate for  $P(y | \text{the})$ ,

$$\frac{1}{2} = \frac{e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}}}}{e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{cat}}} + e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}}}}$$

We get,

$$e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{cat}}} = e^{\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}}}$$

We get,

$$\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{cat}} \approx \mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}} \approx 0$$

Repeating these steps, for  $P(y | a)$ ,  $P(y | \text{cat})$ , and  $P(y | \text{dog})$  we get,

$$\mathbf{v}_{\text{cat}} \cdot \mathbf{c}_{\text{the}} = 0 \text{ and } \mathbf{v}_{\text{cat}} \cdot \mathbf{c}_a = 0$$

and,

$$\mathbf{v}_{\text{dog}} \cdot \mathbf{c}_{\text{the}} = 0 \text{ and } \mathbf{v}_{\text{dog}} \cdot \mathbf{c}_a = 0$$

and,

$$\mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{dog}} = 0 \text{ and } \mathbf{v}_{\text{the}} \cdot \mathbf{c}_{\text{cat}} = 0$$

and,

$$\mathbf{v}_a \cdot \mathbf{c}_{\text{dog}} = 0 \text{ and } \mathbf{v}_a \cdot \mathbf{c}_{\text{cat}} = 0$$

Since,  $\mathbf{v}_{\text{cat}} = (1, 0)$  and  $\mathbf{v}_{\text{dog}} = (1, 0)$ ,

and

$$\mathbf{v}_{\text{the}} = (0, 1) \text{ and } \mathbf{v}_a = (0, 1)$$

The solutions that satisfy the above equations give us the following final word vectors:

- $\mathbf{v}_{\text{the}} = (0, 1)$
- $\mathbf{v}_{\text{dog}} = (1, 0)$
- $\mathbf{v}_{\text{cat}} = (1, 0)$
- $\mathbf{v}_a = (0, 1)$

and context vectors,

- $\mathbf{c}_{\text{the}} = (0, 1)$
- $\mathbf{c}_{\text{dog}} = (1, 0)$
- $\mathbf{c}_{\text{cat}} = (1, 0)$
- $\mathbf{c}_a = (0, 1)$

## References

1. The code uses supplementary Python functions and classes supplied in the PA1 zip folder.
2. The code uses the BPE Algorithm referenced in the lecture slides.
3. The (LaTeX) formatting for this report was prepared with the help of ChatGPT.
4. The report content was revised (for better grammar and sentence restructuring) using Grammarly.